

Univariate Linear Regression

Task 2: Load the Data and Libraries

```
In [1]: import matplotlib.pyplot as plt
plt.style.use('ggplot')
%matplotlib inline
```

```
In [4]: import numpy as np
import pandas as pd
import seaborn as sns
plt.rcParams['figure.figsize'] = (12, 8)
```

```
In [5]: data = pd.read_csv('bike_sharing_data.txt')
data.head()
```

```
Out[5]:
```

	Population	Profit
0	6.1101	17.5920
1	5.5277	9.1302
2	8.5186	13.6620
3	7.0032	11.8540
4	5.8598	6.8233


```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 97 entries, 0 to 96
Data columns (total 2 columns):
Population    97 non-null float64
Profit        97 non-null float64
dtypes: float64(2)
memory usage: 1.6 KB
```

Task 3: Visualize the Data

```
In [8]: ax = sns.scatterplot(x='Population',y='Profit',data=data)
ax.set_title('profit in $10000 v/s city population in $10000')
```

```
Out[8]: Text(0.5, 1.0, 'profit in $10000 v/s city population in $10000')
```



Task 4: Compute the Cost $J(\theta)$

The objective of linear regression is to minimize the cost function

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

where $h_{\theta}(x)$ is the hypothesis and given by the linear model

$$h_{\theta}(x) = \theta^T x = \theta_0 + \theta_1 x_1$$

```
In [9]: def cost_fun(X, y, theta):
m = len(y)
y_pred = X.dot(theta)
error = (y_pred - y)**2
return 1/(2*m) * np.sum(error)
```

```
In [10]: m = data.Population.values.size
X = np.append(np.ones((m, 1)), data.Population.values.reshape(m, 1),axis = 1)
y = data.Profit.values.reshape(m, 1)
theta = np.zeros((2,1))
cost_fun(X,y,theta)
```

```
Out[10]: 32.072733877455676
```

Task 5: Gradient Descent

Minimize the cost function $J(\theta)$ by updating the below equation and repeat until convergence

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \text{ (simultaneously update } \theta_j \text{ for all } j)$$

```
In [11]: def gradient(X,y,theta, alpha,iterations):
m = len(y)
costs = []
for i in range(iterations):
y_pred = X.dot(theta)
error = np.dot(X.transpose(),(y_pred - y))
theta -= alpha * 1/m *error
costs.append(cost_fun(X,y,theta))
return theta, costs
```

```
In [16]: theta, costs = gradient(X, y, theta, alpha = 0.01, iterations=2000)
print('h(x) = {} + {}x1'.format(str(round(theta[0],2)) , str(round(theta[1],2))))
```

```
h(x) = -3.9 + 1.19x1
```


```
In [ ]:
```

Task 6: Visualising the Cost Function $J(\theta)$

```
In [17]: from mpl_toolkits.mplot3d import Axes3D
```

```
In [19]: theta_0 = np.linspace(-10,10,100)
theta_1 = np.linspace(-1,4,100)
cost_values = np.zeros((len(theta_0),len(theta_1)))
for i in range(len(theta_0)):
for j in range(len(theta_1)):
t = np.array([theta_0[i], theta_1[j]])
cost_values[i,j] = cost_fun(X,y,t)
```

```
In [22]: fig = plt.figure(figsize = (12, 8))
ax = fig.gca(projection='3d')
surf = ax.plot_surface(theta_0, theta_1, cost_values, cmap = 'viridis')
fig.colorbar(surf, shrink=0.5, aspect=5)
plt.xlabel('$\Theta_0$')
plt.ylabel('$\Theta_1$')
ax.set_zlabel('$J(\Theta)$')
ax.view_init(30,30)
plt.show()
```

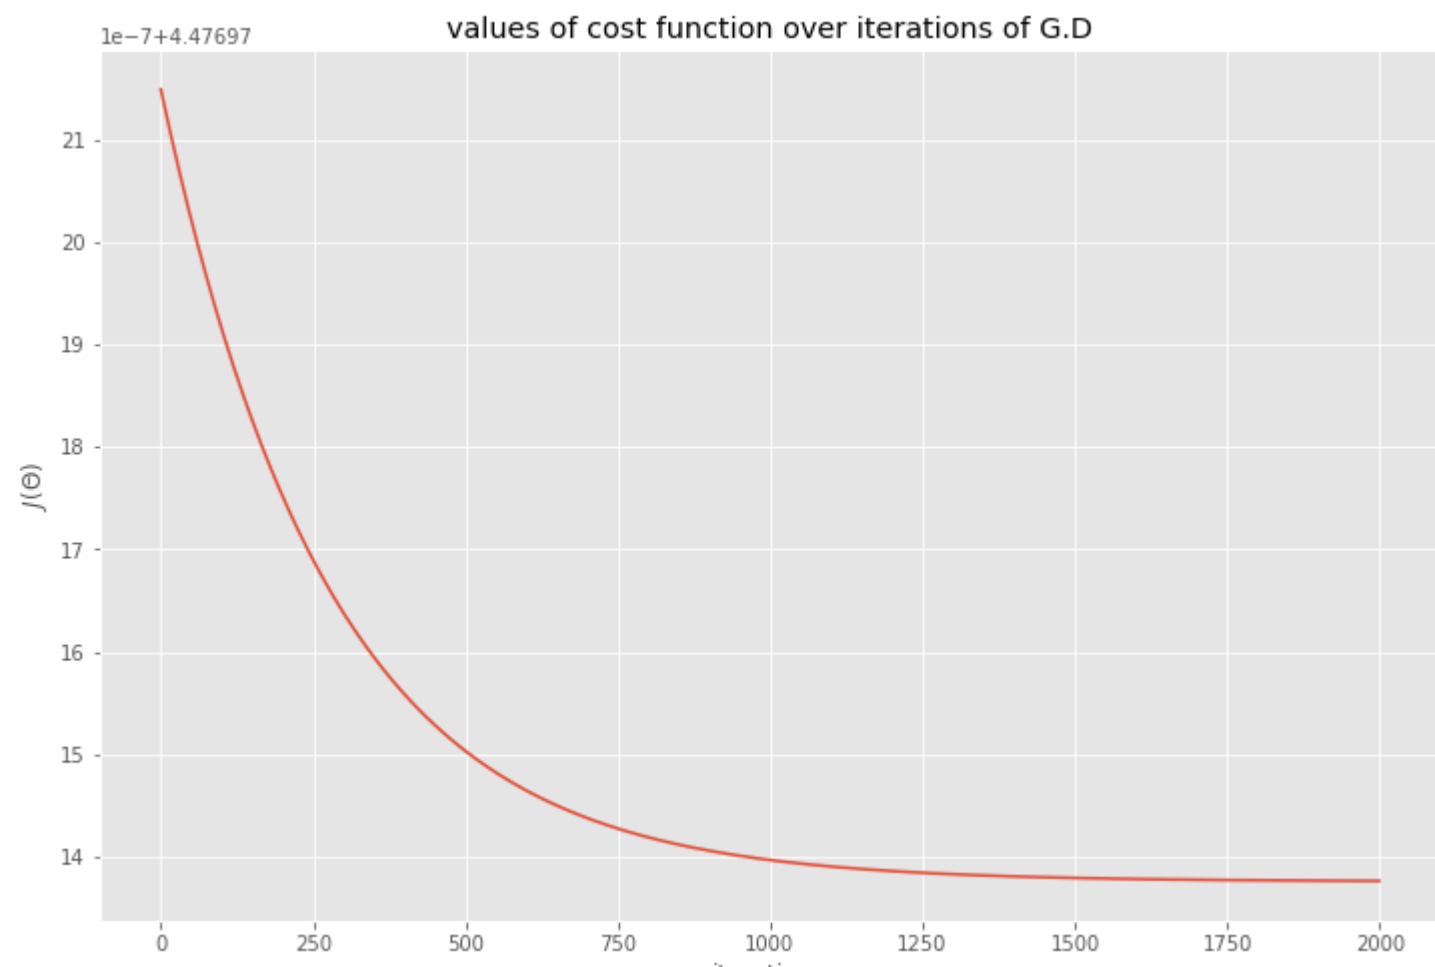


Task 7: Plotting the Convergence

Plot $J(\theta)$ against the number of iterations of gradient descent:

```
In [23]: plt.plot(costs)
plt.xlabel('iterations')
plt.ylabel('$J(\Theta)$')
plt.title('values of cost function over iterations of G.D')
```

```
Out[23]: Text(0.5, 1.0, 'values of cost function over iterations of G.D')
```



Task 8: Training Data with Linear Regression Fit

```
In [24]: theta.shape
```

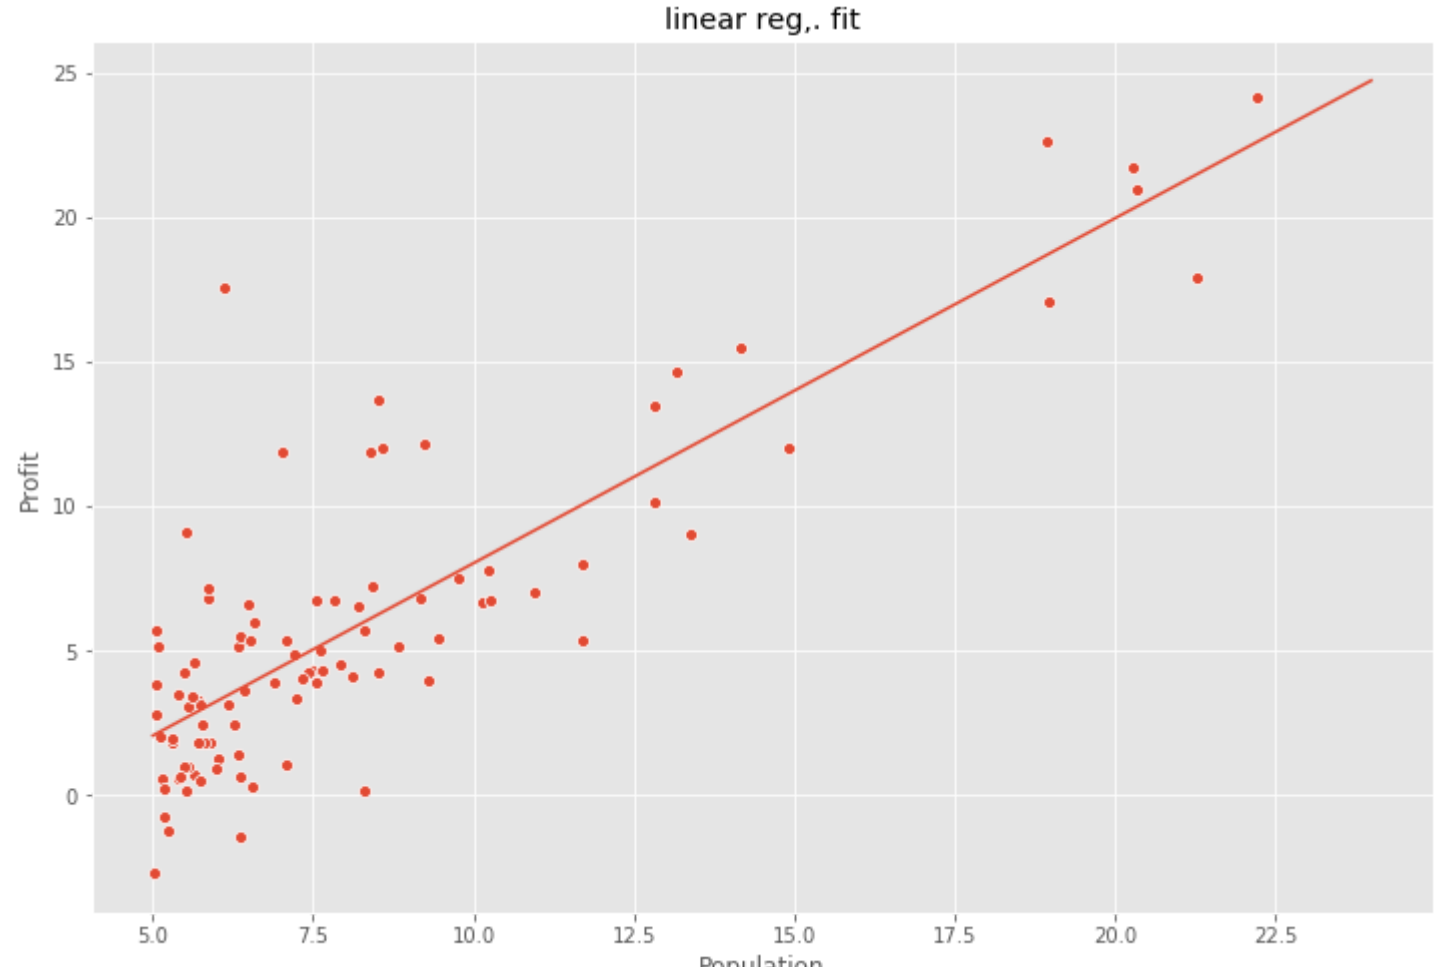
```
Out[24]: (2, 1)
```

```
In [25]: theta
```

```
Out[25]: array([[ -3.89570181],
[  1.1930257 ]])
```

```
In [30]: theta = np.squeeze(theta)
sns.scatterplot(x='Population',y='Profit',data=data)
x_value = [x for x in range(5,25)]
y_value = [(x * theta[1] + theta[0]) for x in x_value]
sns.lineplot(x_value,y_value)
plt.xlabel("Population")
plt.ylabel("Profit")
plt.title("linear reg,. fit")
```

```
Out[30]: Text(0.5, 1.0, 'linear reg,. fit')
```



Task 9: Inference using the optimized θ values

$$h_{\theta}(x) = \theta^T x$$

```
In [31]: def predict(x, theta):
y_pred = np.dot(theta.transpose(),x)
return y_pred
```

```
In [32]: y_pred1 = predict(np.array([1,4]), theta)*1000
print('for population of 40k people, the model predicts a profit of $'+ str(round(y_pred1, 0)))

for population of 40k people, the model predicts a profit of $876.0
```

```
In [33]: y_pred2 = predict(np.array([1,8.3]), theta) *1000
print('for population of 83k people, the model predicts a profit of $'+ str(round(y_pred2, 0)))

for population of 83k people, the model predicts a profit of $6006.0
```